sarvam

AI FOR ALL, FROM INDIA

GOVERNMENT OF CHHATTISGARH

A Full-Stack Sovereign AI Programme for Government and for Citizens

A Grand Vision, Delivered in Modular Phases

Theme 1 — AI for Government (G2G) · Theme 2 — AI for Citizen Services (G2C)

Prepared by Sarvam · Axonwise Private Limited, Bengaluru

July 2026 | Draft for Discussion

Contents

1 The Grand Vision

Chhattisgarh — home to over 3 crore citizens across 33 districts, a large tribal population, and vast tracts where the nearest government office is a long journey away — has both the need and the opportunity to become one of India's leading states in applied artificial intelligence. In a state this diverse, even modest improvements in how government serves, governs, and operates compound across crores of people, in Hindi and in Chhattisgarhi, and in every arm of administration from Raipur to the remotest gram panchayat.

The vision is a state that reasons in its own language: where every official has a sovereign AI assistant that drafts, translates, analyses, and recalls the department's history on demand; and where every citizen can reach a government service, in Chhattisgarhi or Hindi, through a simple phone call or a WhatsApp message, around the clock, without needing literacy in English or a visit to a government office. The same digital public infrastructure approach that took Aadhaar, UPI, and DigiLocker to hundreds of millions can now be extended to AI — this time in the citizen's mother tongue.

> **The programme, in one sentence**
> The Government sets the priorities and supplies the use cases, a System Integrator assembles and integrates the solution, sovereign infrastructure powers it, and Sarvam supplies the sovereign AI models, the agentic platform, and the applications — organised so the State can start small, prove value, and scale as budgets are released.

Why this proposal is structured the way it is

This document is built on two ideas that reflect exactly what the Government of Chhattisgarh asked for. First, a grand vision — a single, coherent picture of an AI-first state that leadership can align behind. Second, a modular architecture — the same vision decomposed into independent, self-contained modules, each of which delivers value on its own and can be switched on in a phased manner, matched to the quantum of funds the State is able to release at any given time. No module depends on the whole programme being funded up front; each is a complete, demonstrable step.

Two themes

Theme	Audience	What it delivers
Theme 1 — AI for Government	Government to Government (G2G)	A sovereign AI workspace for every officer (Indus), plus knowledge and data-verification platforms that turn the State's own records into searchable, conversational institutional memory.
Theme 2 — AI for Citizens	Government to Citizen (G2C)	Population-scale citizen services in Hindi and Chhattisgarhi — voice and WhatsApp agents for grievance redressal, beneficiary outreach, maternal health, scholarships, and land records.

Both themes run on one shared sovereign foundation — the same Indic foundation models, the same agentic platform, the same security and governance controls — so investment in one theme strengthens the other, and the State never pays twice for the same capability. The remainder of this note sets out that foundation (Section 2), the modular phasing tied to budget release (Section 3), and then each theme in detail (Sections 4 and 5).

2 The Sovereign Foundation: One Stack, Every Use Case

Before either theme, it helps to see the common stack both run on. A full-stack sovereign AI system is built from distinct layers, each with a clear function, stacking from sovereign infrastructure at the base to the citizen- and officer-facing applications at the top, with data, security, and governance spanning the whole. No single layer delivers an AI-first state on its own; the value comes from the layers working together — and from the fact that Sarvam builds and owns the layers that matter most: the models, the agentic platform, and the applications.

Layer	Role in the stack	Provided by
Infrastructure & compute	Sovereign cloud, GPU, networking, storage — hosted within Indian territory.	State / SI, or bundled by Sarvam
Foundational models	Language, speech, vision, and translation models pre-trained for Indian languages and context.	Sarvam (OEM)
Agentic platform	Plans and executes multi-step workflows, routes each request to the right-sized model, enforces guardrails.	Sarvam (OEM)
Data & integration	Pipelines, lakehouse, vector databases, connectors to departmental systems, LLMOps, identity.	Sarvam + SI
Application layer	Officer- and citizen-facing products that deliver outcomes in the citizen's own language.	Sarvam (OEM)
Security & governance	Encryption, access control, audit trails, responsible-AI controls — embedded at every layer.	Embedded throughout

Pillar 1 — Sovereign and Foundational Models

Sarvam has built the full model stack in India: data curation, tokenizer, architecture, pre-training, post-training, inference, speech, vision, and translation. All models are developed and trained in India and can run entirely within Indian territory. The weights are owned by Sarvam, and the flagship reasoning models (Sarvam 30B and Sarvam 105B) are open-sourced under Apache 2.0, so the State can host them within its own boundary rather than rent them over an API it does not control. A single model size does not fit every government workload, so Sarvam offers a family sized to the task.

Tier	Parameter range	Primary use
Lightweight / on-device	3B	Summarisation, translation, low-latency and low-cost tasks; deployable on edge devices.
Reasoning / flagship	105B	Complex reasoning, agentic workflows, synthetic-data generation, coding with think mode.
General-purpose engine	500B-class (in pipeline)	Long-horizon reasoning, multi-step tool use, agentic coordination across agents.

The stack is multimodal from the start — speech, vision, and translation models make the language models useful across the real working surfaces of officers and citizens.

Model	Type	Key capabilities
Saaras V3	Speech-to-Text	Streaming recognition across 22 Indian languages; code-mixed support; automatic language detection and speaker diarization; sub-250ms latency.

Model	Type	Key capabilities
Bulbul V3	Text-to-Speech	Natural, expressive speech across 11 Indian languages; 6+ voices per language; control over pitch, speed, and style.
Sarvam Vision	Vision-language (OCR)	3B model for document digitisation and OCR in Indian scripts; layout-preserving extraction of text, tables, and forms.
Sarvam Translate / Mayura	Translation	Translation across 22 Indian languages, including colloquial and code-mixed text, preserving formatting.

Pillar 2 — Agentic Platform Layer

Models supply reasoning; the agentic platform turns that reasoning into reliable, accountable work. It plans and coordinates multi-step workflows, routes each query to the right-sized model, manages memory across sessions, calls departmental systems, and enforces guardrails and human approvals before any consequential action — with a durable runtime and a trace store that make every government workflow recoverable and auditable. Crucially, it is model-agnostic and centrally governed: every agent deployed across the State, in any department, routes through one platform, so logs, guardrails, and policy updates are applied once, centrally, rather than rebuilt in each department.

Pillar 3 — Application Layer

The application layer turns the models and the platform into surfaces officers and citizens actually use. It is here that the two themes of this proposal live — Indus and the Knowledge and Data platforms for government (Theme 1), and Samvaad voice/WhatsApp agents and citizen platforms (Theme 2) — all drawing on the same Pillar 1 and Pillar 2 foundation.

3 Modular Architecture and Budget-Linked Phasing

The defining design principle of this programme is that the grand vision is delivered as independent modules. Each module is a complete capability that stands on its own, produces a visible outcome, and can be procured and switched on when the State is ready to release the corresponding funds. The State is never asked to commit the full programme budget up front. Instead, it climbs a value ladder — each rung inexpensive to start, each proving return before the next is funded.

How the modularity works

- Independent modules. Each module — an officer workspace, a grievance-redressal voice line, a land-records digitisation drive — is deployable without the others. A department can buy exactly one and derive full value from it.
- Shared foundation, paid once. Every module reuses the same models, agentic platform, and security controls. Funding the first module stands up the foundation; later modules cost less because that foundation already exists.
- Budget-linked activation. Modules map to fund tranches. A small pilot budget activates a Phase-1 module; a larger release in the next quarter or financial year activates the next. Nothing stalls waiting for a single large sanction.
- Value ladder maps to the buyer journey. Low-ticket pilots (search, Q&A, a single helpline) get a foot in the door; dashboards and analytics win over Secretaries; collaboration and agentic workflows make the platform sticky and hard to displace.

Indicative phasing tied to fund release

The phases below are illustrative. Any phase can be scoped up or down, and a department may sit at a different phase from its neighbour under the same master contract. The point is that each phase is a self-funding unit of value.

Phase	Fund release	Modules activated	Outcome
Phase 1 — Foundation & Quick Wins	Smallest tranche / pilot budget	Sovereign foundation stood up; Indus base workspace for 1–2 priority departments; one citizen voice/WhatsApp helpline; one document-digitisation pilot.	Demonstrable value within weeks; leadership sees working AI on the State's own data and citizens.
Phase 2 — Departmental Expansion	Mid-size release (next quarter / FY)	Indus rolled to more departments; department knowledge ingestion and dashboards; grievance redressal and beneficiary outreach at scale; data-verification platform for land / scholarships.	AI embedded in daily operations of multiple departments; Secretaries see analytics and MIS in plain language.
Phase 3 — State-wide Scale	Programme-scale funding	State-wide Indus fabric; citizen general-purpose assistant; agentic workflows; RBAC, collaboration, eOffice-style integration; full G2C service coverage.	A common AI fabric across all departments and a single front door for citizens; capacity transferred to State teams.

Why this protects the State

> The State pays in proportion to value already proven, retains a clean exit at every phase (the flagship models are open-source and run inside the State's boundary), and is never locked into a large irreversible commitment. Each rupee released buys a working, standalone capability — not a promise contingent on the next sanction.

4 AI for Government — Government to Government (G2G)

Officers spend a large part of their day moving information between systems — drafting notings in one file, looking up history in another, writing replies in Word, copying figures from a spreadsheet, searching for old circulars, preparing briefs from scattered sources. Theme 1 gives the Government of Chhattisgarh a sovereign AI layer for its own people: a workspace for every officer, and the platforms that turn the State's own records and data into searchable, conversational institutional memory. Three modules make up this theme; each can be activated independently.

Module	What it is	Typical first phase
4.1 Indus — AI Workspace for Officers	A sovereign AI productivity workspace every officer logs into.	Base subscription, 1–2 departments
4.2 Knowledge Platform	Turns a department's document archives into voice/text Q&A with citations.	Ingest + search for one department
4.3 Data Verification Platform	Bulk-digitises handwritten records and verifies them against canonical sources.	Land records or scholarship pilot

4.1 Indus — A Sovereign AI Workspace for Every Officer

Indus brings an officer's work into a single sovereign workspace, hosted in India and operated under Indian government rules. Inside the workspace, an AI assistant reads the officer's files, recalls the department's history, drafts documents, analyses data, and produces ready-to-share material on demand. Every officer in the Government of Chhattisgarh gets sovereign, secure access to AI — built and operated in India by Sarvam.

What the AI agents in Indus can do

- Research, drafting, and review. Deep research, drafting, translation, summarisation, comparison, and decision support across English, Hindi, Chhattisgarhi, and other Indian languages.
- Document understanding. Read and analyse PDFs, Word files, spreadsheets, scanned papers, and images, with high-accuracy OCR across Indian scripts.
- Ready-to-share artifacts. Generate file notings, briefs, presentations, reports, and dashboards in the formats government uses.
- Data analysis without code. Upload an MIS download, scheme file, or budget statement and ask questions in plain language — get charts, comparisons, and exception analyses.
- Pre-built workflow skills. Common workflows like drafting a file noting or responding to RTI applications are built in as ready-to-use skills.
- Reasoning on the department's own knowledge. Where the department brings its files, circulars, and records into Indus, the agent answers from that knowledge with citations to source documents.
- Multi-channel access. Web, desktop, mobile, WhatsApp, voice, or embedded into government portals.

Three ways officers use Indus

Indus is useful from Day 1 and becomes more useful as more context, files, and tools connect to it. There are three ways to use it.

Way 01 · Your Personal Workspace. One screen to write notes, upload files, ask questions, get summaries, draft replies, prepare for meetings, and produce presentations. Once connected to email, calendar, and eOffice, it generates a personalised morning brief showing pending files, today's meetings with prep materials, and threads awaiting reply. This layer saves the most time per officer, per day.

> An everyday example
> Secretary, IT · 7:15 AM, on the way to office. *The Secretary opens Indus on the phone. The morning brief shows four files pending noting in eOffice, three email threads awaiting reply, and today's calendar with prep briefs ready for two meetings. He skims them in the car and walks in already prepared.*

Way 02 · Working as a Team. For a scheme review, a procurement, or a policy exercise, officers create a project room where files, decisions, action items, meeting minutes, and the AI agents working alongside the team stay in one place. When someone is transferred or a new colleague joins, the project's history is intact.

> An everyday example
> Procurement team, PWD · ICT system tender. *An RFP Drafting & Compliance skill drafts the tender through interactive Q&A — asking about experience thresholds, turnover, and evaluation weightage — and assembles a complete, GFR-compliant document. When bids arrive, the agent generates a compliance matrix against every criterion, a comparative scorecard, and flags missing documents and abnormally low bids. The evaluation report writes itself.*

Way 03 · Tapping into Department Memory. Every department holds two kinds of memory — the historical record built up over decades (files, circulars, committee reports, scheme guidelines, court orders) and the live operational pulse (MIS systems, beneficiary databases, field reports). Indus connects to both, and the AI answers from the department's own institutional memory, in plain language, with citations to source.

> An everyday example
>
> Additional Collector, Bastar · in the car, on the way to a field visit. *On the way to inspect PMAY-G construction in two villages, she asks Indus in Hindi how many houses have been sanctioned and completed under PMAY-G in the villages she is visiting today. Indus answers with numbers, the names of beneficiaries whose construction is stalled, and the last reported reason from the block office.*

The principle across all three ways is the same: officers spend less time searching, compiling, and reformatting — and more of the day on work that needs their judgment. Indus does not replace the officer; it clears away the work around the officer.

Capability areas — what Indus can do today

Capability area	What becomes possible	Illustrative use cases
Personal AI productivity	A secure AI workspace for research, drafting, analysis, review, and artifact creation.	Notes, briefs, presentations, reports, meeting prep, first drafts of replies.
Document intelligence	Read and analyse any document — PDFs, Word, spreadsheets, scanned or handwritten notes.	Compare policy versions, extract tables, summarise long files, digitise legacy archives.
Multilingual interaction	Full text and voice across 22 Indian languages and English; code-mixed speech native.	Hindi and Chhattisgarhi file notings, GR drafting, bilingual outputs in parallel.
Pre-built workflow skills	Common government workflows packaged as ready-to-use skills from Day 1.	Noting & Drafting, RTI Response — trained on the relevant rules, formats, precedents.
Project workspaces	Multi-officer rooms with files, decisions, meeting memory, and AI agents — persisting across transfers.	Budget reviews, policy drafting, scheme monitoring, procurement evaluation rooms.
Department knowledge	Connect Indus to historical archives, live MIS, and operational databases.	Query decades of policy files, track litigation, ask field-data questions in plain language.
Security & sovereignty	Identity, permissions, auditability, source provenance, full data residency in India.	Role-based access by designation, tamper-evident audit trails, citations, optional air-gap.

How Sarvam delivers Indus — two tracks

Track One · The Base Platform. Sarvam operates the base Indus platform as a managed AI productivity layer for the State. The base subscription includes everything needed to be productive from Day 1 — the home AI agent every officer gets, pre-built skills (Noting & Drafting, RTI Response), document intelligence, artifact creation, multilingual voice interaction, project workspaces, secure data analysis, multi-channel access, standard integrations where APIs exist (eOffice, NIC email, calendar), ingestion of Chhattisgarh-specific public knowledge (the State's Manual of Office Procedure and procurement rules), and administrator controls with tamper-evident audit logs. Departments do no engineering; officers log in and use begins the same week. Priced per officer, per month, at the same rate across departments.

Track Two · Department Transformation. Sarvam places Forward Deployed Software Engineers directly inside a department to build the agents and integrations it needs — ingesting historical files into institutional memory, building custom workflow agents (an Assembly Q&A agent grounded in past replies, a Cabinet Note agent grounded in the State's decision history, an audit-response agent), integrating live MIS and scheme systems, and, for sensitive departments such as Home and Police, deploying air-gapped on the department's own hardware. Each engagement is scoped and priced jointly with the department.

> Indus is infrastructure, not just an app
>
> Over time, departments deploy custom agents of their own — from scheme monitoring to citizen grievance triage to industry consultation — without each having to build its own AI stack. That is the modularity of Theme 1: one base platform, then as many department-specific transformation modules as the State chooses to fund, when it chooses to fund them.

4.2 Knowledge Platform — Institutional Memory, Made Conversational

At its core, this platform takes a messy folder of government documents, normalises them using Sarvam's vision and translation models, organises them into a structured ontology, and exposes that structure through search, voice-based Q&A, and — in later phases — dashboards and document creation. It is the State-government surface of Sarvam's Knowledge Fabric, and it is deliberately built so a non-technical GTM person can upload a department's circulars and demo policy search within an hour.

Use cases

- Legal case assistant — helps legal and administrative departments classify, manage, and act on their case backlog.
- Frontline worker (e.g. ASHA) assistant — puts institutional knowledge into the hands of field workers.
- Citizen assistant for public-services search — natural-language navigation of State-specific public services.
- Internal productivity co-pilot and MDR / FBMDR analysis.

Phased roadmap (module-by-module)

Phase	What ships	Value unlocked
Phase 1 — Ingest, Structure, Search	Ingestion pipeline (vision normalisation + Indic translation), interactive ontology creation, and voice-based Q&A over documents.	Demo-ready, GTM-ready foundation. Frontline workers get value from day one.
Phase 2 — Dashboards, Doc Creation, RBAC	Dynamic dashboards from metadata filters, template-driven document drafting, and role-based access control.	Wins over Directors and Secretaries — the decision-makers who greenlight projects.
Phase 3 — Collaboration & Sharing	eOffice-style in-app document sharing, routing, and organisational workflows.	Becomes the single working environment — sticky and hard to displace.
Phase 4 — Agentic Workflows	Auto-flagging for approval, intelligent routing, automatic prioritisation of to-dos.	Agentic inbox-outbox management, architected for now, activated as reliability matures.

> Design principles
>
> GTM-first setup — a working Q&A demo within an hour of uploading a folder. Indic-native — translation and vision normalisation are in the core pipeline, not add-ons. Value ladder maps to the buyer journey — Phase 1 opens the door, Phase 2 wins Secretaries, Phase 3+ makes it permanent.

4.3 Data Verification Platform — From Handwritten Records to Verified Data

This platform takes scanned, usually handwritten, documents as input, along with a few samples showing the fields to extract. From this it supports bulk ingestion, extraction, and human review — and, in later phases, rule-based verification against canonical sources with automated next steps. It reuses the Knowledge Platform's ingestion and normalisation pipeline, so the two are two surfaces on shared infrastructure.

Use cases

- Land-record digitisation and mutation processing.
- Scholarship and other welfare disbursals.
- Beneficiary verification for maternal-health benefits, and general beneficiary management.

Phased roadmap

Phase	What ships
Phase 1 — Ingest, Extract, Review	Sample-based schema definition (annotate 3–5 documents), bulk ingestion and extraction through vision + Indic normalisation, and a human-in-the-loop review UI.
Phase 2 — Verification Against Canonical Sources	Register reference datasets (revenue-court records, MIS databases, validation APIs); a configurable match-rule engine produces per-field verdicts (confirmed / mismatch / unverifiable), surfaced in a green-amber-red discrepancy dashboard.
Phase 3 — Notifications	Rule-based internal notifications to officials; optional citizen-facing templated status updates (e.g. scholarship or land-mutation status).
Phase 4 — Bulk Re-Processing	Automatic re-verification when canonical sources update (new geospatial survey, refreshed revenue DB), with delta reports on records whose status changed.
Phase 5 — Action Recommendations at Upload	Live pre-verification triage — missing mandatory fields, out-of-range values, in-batch duplicates flagged as documents are extracted.

> Design principles
>
> Human-in-the-loop by default — every extraction and verification result is reviewable. Audit trail — every extraction, correction, match, and action is logged with who, when, and what changed. Indic-native — vision and translation models are core pipeline dependencies.

THEME 2

5 AI for Citizen Services — Government to Citizen (G2C)

Theme 2 is the citizen-facing half of the programme: population-scale services delivered in the citizen's own language and dialect, through the channels citizens already use — a phone call or a WhatsApp message. It is powered by Samvaad, Sarvam's conversational-agents platform, together with the beneficiary and verification platforms behind it. The use cases below are drawn from live deployment conversations Sarvam is already having with State governments, clustered so that Chhattisgarh can activate whichever citizen service is the priority, when the budget for it is released.

Samvaad — the citizen-facing delivery layer

Attribute	Specification
Channels	Voice (inbound and outbound), WhatsApp, and web chat from a single unified platform.
Languages	22 Indian languages with dialect and code-switching support — Hindi and Chhattisgarhi, voice-in and voice-out for low-literacy citizens.
Latency	Time-to-first-voice below 2 seconds (p95); sub-500ms end-to-end voice response.
Concurrency	Population-scale; demonstrated 1,000+ concurrent sessions across 10+ Indian languages with active tool calls.
Orchestration	Multi-agent execution across departmental workflows, with human handoff for anything sensitive.
Time-to-launch	A new agent in under 24 hours via pre-built government templates.

The five citizen-service modules that follow are each independent. Each is a candidate for a Phase-1 quick win on its own, and each deepens as budget allows.

5.1 Citizen Grievance Redressal

Municipalities and departments receive a large volume of citizen queries and complaints through helplines — many of them repetitive status requests or standard service issues that today require manual call handling and ticket creation. This leads to long wait times, inconsistent categorisation, poor data quality, and limited visibility into complaint trends. A Samvaad grievance agent answers the citizen directly, classifies intent, captures structured data, raises work orders, and escalates to a human only when the matter genuinely needs one.

What the agent handles

- Information requests. Birth/death certificate status, property-tax or water-bill queries, building-permit status, scheme eligibility — retrieved and resolved without human intervention.
- Service requests. Garbage not collected, streetlight out, water leakage, drain blockage, potholes — the agent captures location, auto-generates a work order, assigns the department, and gives the citizen a ticket number and ETA.
- Inspection-based complaints. Illegal construction, encroachment, illegal dumping — the agent collects details, photos, and location and creates an inspection case for human review.
- Human escalation. Disputes, corruption allegations, legal notices, threats to public safety, or repeatedly unresolved complaints are captured and transferred to the right human team.

> Feature set for deployment
>
> Structured intake (intent classification into municipal service categories, field extraction, standardised grievance codes); standardised disposition capture (resolution categories and reason codes); workflow automation (auto work-order and inspection-case generation, SLA tagging); and case management (unique tickets, status retrieval, automatic escalation on SLA breach).

Requested across State engagements by welfare and municipal departments; Chhattisgarh can pilot on a single service line (e.g. a municipal helpline or a scheme grievance line) and expand.

5.2 Maternal Health — Proactive Outreach and Advisory

A voice-and-WhatsApp companion that accompanies an expectant mother through her pregnancy journey, triggering timely AI calls at the key clinical milestones and answering her questions on WhatsApp in between. Registered against the maternal journey at EDD calculation, the agent proactively calls at defined stages, checks progress, resolves gaps live, and follows up with WhatsApp guidance.

The proactive call journey

Trigger	Timing	What the agent does
First ANC checkup	10–12 weeks	Confirms whether the first checkup is booked; if not, books it or shares the nearest facility, then sends appointment details, directions, and a document checklist on WhatsApp.
Iron tablet & nutrition adherence	20–24 weeks	Checks adherence; if the mother has stopped, identifies why (side effects, forgetting, tablets finished) and responds — side-effect guidance, daily reminders, or refill-facility information.
Delivery preparedness	34–36 weeks	Confirms hospital selection, transport, emergency contact, and a ready hospital bag; fills any gap live and sends a delivery checklist.

Researched supporting applications

- Resource Mobilisation Agent — identifies weekly resource requirements per area from the eKavach database and calls PHCs to ensure preparedness.
- Stockout Monitoring Agent — ensures life-saving drugs are physically present at peripheral facilities.
- Beneficiary Verification + District Dashboard — verifies that services the system claims to deliver actually reach pregnant women.
- MDR Analysis Pipeline — converts maternal death review records from filing-cabinet artifacts into systemic learning.
- Family Decision-Maker Engagement Agent — engages the household decision-maker, especially during emergencies.

The advisory bot and the automatic CMS workflow agent can be deployed independently, and the outreach runs entirely in Hindi and Chhattisgarhi by voice.

5.3 Beneficiary Management and Auto-Enrolment

This platform links to existing State databases to validate citizen eligibility across open government schemes and triggers consent calls for auto-enrolment — moving the State from a system where citizens must find schemes to one where schemes find eligible citizens.

Use cases

- Complete profile creation after triaging user data from multiple sources.
- Automatic verification against stated scheme criteria.
- Auto-enrolment across State schemes.

Required features

1. Data connection. Connect to State databases holding key user-eligibility fields.
2. Eligibility criteria. Department-wise definitions of the key eligibility rules.
3. Past-beneficiary linkage. Link to the database of past beneficiaries per scheme.
4. Auto-verification workflows. Trigger consent collection or data re-verification through Samvaad voice agents (Phase 2).
5. Final beneficiary list. The latest-period list forwarded to the relevant department for disbursal.

5.4 Land Records and Scholarship Digitisation

The Data Verification Platform (Section 4.3) has a direct citizen-facing surface. For land records, the State can digitise legacy records, mutation orders, and RoR documents — including handwritten annotations in regional scripts — and then triangulate the extracted data against canonical sources such as geospatial and revenue records. For scholarships and welfare, it supports end-to-end collection and digitisation of student data at enrolment, automatic verification, annual re-verification, and disbursements linked to the student profile — with citizens receiving templated status updates on their application.

> Why it belongs in both themes
>
> The same platform serves an internal government need (Theme 1: clean, verified data) and a citizen outcome (Theme 2: faster, transparent land mutation and scholarship disbursal). Fund it once, and both themes benefit — a concrete example of the shared-foundation economics that run through this proposal.

5.5 Citizen Helpline and Public-Services Navigator

A 24×7 multilingual inbound helpline and a general public-services navigator give citizens a single front door to State schemes — eligibility, application, status, and grievance routing — in Hindi and Chhattisgarhi, by voice or WhatsApp. This is the natural convergence point of the theme: as grievance redressal, beneficiary outreach, and scheme databases come online, the citizen navigator ties them together into one conversational entry point, so a citizen need not know which department owns their problem.

- Citizen helplines. 24×7 multilingual support for welfare schemes, land records, grievance registration, and service-status tracking.
- Proactive outreach. Outbound campaigns to verify beneficiary data, confirm scheme eligibility, and collect feedback at scale.
- Sector companions. Agricultural advisory (crop advice, mandi prices, weather alerts) and health outreach (immunisation reminders, OPD scheduling) through existing frontline networks.

6 Sovereignty, Deployment, and Governance

Sovereignty is the organising principle of this programme, not a feature added at the end. For the State it means that sensitive citizen data stays under State control, that the intelligence serving citizens is owned rather than rented, and that the State is never locked into a capability it cannot operate independently.

Four dimensions of sovereignty

- Model sovereignty. Models developed and trained in India; flagship models (Sarvam 30B and 105B) open-sourced under Apache 2.0, hostable within the State's boundary, with weights, pre-training data, and methodology available for government audit.
- Data sovereignty and residency. All data and inference remain within Indian territory — and within the State's own infrastructure for on-premises deployments. AES-256 at rest, TLS 1.2+ in transit; Customer-Managed Encryption Keys so even Sarvam cannot access customer data at rest for enterprise deployments.
- Infrastructure sovereignty. Deployable entirely on-premises within State data centres, including air-gapped, zero-internet configurations, with no dependency on any foreign cloud.
- Operational independence and exit. Because the flagship models are open and the deployment runs inside the State's boundary, the State is not dependent on a single vendor. Custom models trained on the State's data stay inside the State's environment, with weights the State owns and Sarvam does not reuse.

Deployment models

Mode	Description	Best suited for
On-Premises	Entire platform, models, and data within the State's own data centres (NIC / State Data Centre). No data or inference leaves State infrastructure.	Highly sensitive data (revenue, law enforcement, health); air-gap requirements.
Virtual Private Cloud	Dedicated, single-tenant cloud within India with full data isolation; State retains control over retention, access, and audit logs.	Sensitive citizen data needing dedicated infrastructure without owning GPU.
SaaS	Shared, multi-tenant cloud within India. Fastest and most cost-effective; automatic upgrades; strict access controls.	Standard sensitivity; rapid rollout; pilots and quick wins.

Delivery and governance

Sarvam deploys a dedicated team of forward-deployed engineers and delivery leads based in or near Raipur. This team learns departmental workflows, co-designs solutions with government stakeholders, and owns testing and deployment end-to-end — which is how capacity is transferred to the State over time. Sarvam partners with one or more System Integrators for scaled development and integration with individual department systems, keeping deep AI expertise at the core and scalable delivery capacity on the ground. A designated nodal agency (such as the IT Department or Chief Secretary's Office) anchors the programme, a Steering Committee chaired by the Chief Secretary reviews progress, and monthly operational and quarterly strategic reviews track deployment, usage, citizen impact, and financial utilisation.

Compliance and standards

Framework	Status
ISO/IEC 27001:2022	Certified

Framework	Status
SOC 2 Type II	Certified
ISO/IEC 42001 (AI Management)	In progress
DPDP Act 2023	Aligned; certification in progress
MeitY guidelines & CERT-In directives	Aligned — applied across UIDAI, NPCI, and IndiaAI deployments

7 Customer References

The capabilities set out in this proposal are not aspirational — they are live, in production, at population scale. The references below are a representative selection of Sarvam's deployments across the Government of India, sovereign digital-public-infrastructure bodies, and large regulated enterprises. Each entry notes what the engagement establishes as a credential for a government buyer such as the Government of Chhattisgarh — sovereign deployment, population-scale multilingual delivery, security posture, and proven government adoption. Client contact details and commercial terms are available on request under mutual NDA.

Reference deployments at a glance

Client	Engagement	Since	Government credential established
NITI Aayog / Office of the Chief Economic Adviser, Govt. of India	Sovereign reasoning engine for natural-language querying over structured and unstructured government data; committed to draft the Economic Survey 2026.	2025	Sovereign AI decision-support adopted at the highest level of the Union Government.
UIDAI (Unique Identification Authority of India)	Full-stack generative-AI Aadhaar helpdesk in 10 Indian languages with real-time fraud detection, deployed on-premises and air-gapped inside UIDAI infrastructure.	2024	Air-gapped, on-premises sovereign deployment within a strict government security perimeter; DPI-grade security.
NPCI (National Payments Corporation of India)	Conversational speech-based Voice UPI / BBPS interface for feature phones in 22 Indian languages, targeting financially-excluded citizens.	2024	Full 22-language coverage and compliance with national DPI standards for last-mile, low-bandwidth inclusion.
Life Insurance Corporation of India (LIC)	AI voice agents for policy-service intimation and contact-centre augmentation, handling lakhs of citizen conversations.	2024	Population-scale, sensitive citizen interactions for a Government of India enterprise.
SBI Life Insurance	AI voice & WhatsApp platform for 3.5 lakh+ agents in 11 languages; multilingual outbound campaigns reaching 4.5 crore policyholders over telephony and WhatsApp.	2024	Proven 4.5 crore+ beneficiary outreach and multilingual distribution-force enablement.
Tata Capital Limited	Samvaad multilingual voice agents across loan sales, verification, and collections in 11 languages; 20 lakh+ monthly interactions with seamless human handoff.	2024	Enterprise-grade voice AI at scale with audit trails and human-in-the-loop governance.

> **Why these matter for Chhattisgarh**
> Read together, these deployments cover every requirement this programme places on a vendor: a sovereign reasoning engine trusted by the Union Government (NITI Aayog/CEA), an air-gapped on-premises deployment inside a national security perimeter (UIDAI), 22-language last-mile voice on ordinary phones (NPCI), and population-scale citizen outreach (LIC, SBI Life). The same platforms power Theme 1 and Theme 2 of this proposal.

1. NITI Aayog / Office of the Chief Economic Adviser (CEA), Government of India

Field	Detail
Engagement	Sovereign Reasoning Engine for Natural Language Querying, Chain-of-Thought Reasoning & Policy Analytics
What was delivered	Built a Reasoning Engine for NITI Aayog enabling natural language querying over structured and unstructured government data, chain-of-thought reasoning, and interactive data visualisations. The CEA has committed to using this engine to draft the Economic Survey 2026 — marking a breakthrough in AI-assisted government document authoring.
Coverage	New Delhi, India
Live since	2025 · 12+ months (Ongoing)

Government credential: A sovereign reasoning engine trusted to draft the Economic Survey 2026 — proof Sarvam can build sovereign AI decision-support tools adopted at the apex of the Government of India, directly relevant to Chhattisgarh's Officer Copilot and policy-analytics needs.

2. UIDAI (Unique Identification Authority of India)

Field	Detail
Engagement	Air-Gapped Sovereign Generative AI Deployment for Aadhaar Multilingual Support & Fraud Detection
What was delivered	Built a full-stack Generative AI solution for Aadhaar support, deployed in an on-premises air-gapped environment. Solution provides multilingual voice-based Aadhaar helpdesk in 10 Indian languages and performs real-time fraud detection — entirely within UIDAI's secure infrastructure.
Coverage	New Delhi, India
Live since	2024 · 12+ months (Ongoing)

Government credential: An air-gapped, on-premises deployment operating entirely within UIDAI's secure infrastructure — proof Sarvam can meet strict government data-residency and security perimeters, essential for a DPDP-compliant, MeitY-governed sovereign AI substrate.

3. NPCI (National Payments Corporation of India)

Field	Detail
Engagement	Voice UPI — Conversational Speech-Based Payments Interface in 22 Indian Languages
What was delivered	Developed a conversational speech-based UPI and BBPS payments interface for feature phones in partnership with NPCI, targeting 30 Crore financially-excluded Indians. The solution enables UPI transactions in 22 Indian languages for digitally-backward users — fully compliant with NPCI guidelines.
Coverage	Pan India, India
Live since	2024 · 12+ months (Ongoing)

Government credential: Full 22-language, voice-first coverage on ordinary feature phones, compliant with national DPI standards — the closest analogue to reaching print-non-literate and low-bandwidth citizens via telephony, WhatsApp voice, and IVR.

4. Life Insurance Corporation of India (LIC)

Field	Detail
Engagement	AI Voice Agents for Policy Service Intimation & Contact Center Augmentation
What was delivered	Deployed AI voice agents for LIC to intimation customers on policy-related service requests and pending actions. Agents have handled lakhs of conversations to facilitate claims processing and augment LIC's contact center — improving key business metrics at scale.
Coverage	Pan India, India
Live since	2024 · 6+ months (Ongoing)

> Government credential: Population-scale service to a Government of India enterprise's distributed beneficiary base — proof Sarvam can handle sensitive citizen interactions responsibly while augmenting existing contact-centre infrastructure.

5. SBI Life Insurance

Field	Detail
Engagement	AI Voice & WhatsApp Platform for Customer Engagement and Sales Enablement
What was delivered	Deployed AI-powered WhatsApp assistant for 3.5 lakh+ sales agents to answer product queries in 11 languages, generate sales collateral, run premium calculations, and conduct multilingual outbound campaigns (PMJJBY renewals) reaching 4.5 Crore policyholders over telephony and WhatsApp.
Coverage	Pan India, India
Live since	2024 · 12+ months (Ongoing)

> Government credential: Demonstrated 4.5 crore+ beneficiary outreach and multilingual distribution-force enablement across WhatsApp and telephony — validating population-scale, multi-lingual delivery of the kind Chhattisgarh's citizen services require.

6. Tata Capital Limited

Field	Detail
Engagement	Multilingual Conversational AI Agents for Loan Sales, Welcome Calling & Collections
What was delivered	Deployed Samvaad multilingual AI voice agents across Loan Sales, Pre-issuance Verification and EMI Collections in 11 Indian languages. Handles 20 lakh+ monthly interactions with seamless human handoff, delivering 22x+ ROI. Sarvam replaced legacy IVR with natural conversational AI covering full loan lifecycle.
Coverage	Pan India, India
Live since	2024 · 12+ months (Ongoing)

> Government credential: Enterprise-grade Samvaad voice AI at 20 lakh+ monthly interactions across 11 languages with real-time audit trails and human handoff — the same platform that powers the citizen and officer surfaces in this proposal.

8 Why Sarvam

A statewide sovereign AI programme should be anchored by the entity that actually builds the models, not a reseller of someone else's. Sarvam is an AI Original Equipment Manufacturer: it pre-trains its own foundational models from scratch on data it owns, owns the model weights and architecture, and has deployed those models at population scale across government. Owning the weights and training data is not a procurement technicality — it is what makes sovereign control real: a model builder can host weights inside the State's boundary, fine-tune on the State's data without that data leaving, make training data auditable for fairness and safety, and guarantee that a small amount of poisoned data cannot create a hidden backdoor.

Core competency	Sarvam
Owns the model weights	Owns the weights, architecture, and tokenizer of its foundational models; flagship models open-sourced under Apache 2.0 (March 2026).
Owns the training data	Owns the multi-trillion-token pre-training corpora (Sarvam 105B trained from scratch on 12 trillion tokens); training data and methodology available for government audit.
Pre-trains from scratch	Distributed pre-training from scratch in India across 1,000+ GPU clusters, without dependence on a hyperscaler cloud.
Proven deployments	Production deployments across the Ministry of Agriculture, PMO, NCERT, UIDAI, and the National Health Authority, plus State-government programmes.
Indic-first by design	Models built for 22 Indian languages, dialects, and code-mixed speech — not foreign models adapted after the fact.
Security certification	ISO/IEC 27001:2022 and SOC 2 Type II certified; aligned to DPDP, MeitY, and CERT-In requirements.

About Sarvam

Sarvam (Axonwise Private Limited) was founded in August 2023 and is headquartered in Bengaluru. It builds India's full-stack sovereign AI platform, developing every layer — from foundational models to citizen-facing applications — entirely within India. Sarvam was selected by the Government of India under the IndiaAI Mission to build India's sovereign Large Language Model, with active deployments across the Ministry of Agriculture, PMO (Mann Ki Baat dubbing), NCERT (NEP textbook translations), UIDAI, and the National Health Authority, and State-government partnerships active in Tamil Nadu, Odisha, Maharashtra, and beyond. Its founders — Dr. Vivek Raghavan and Dr. Pratyush Kumar — played leading roles in India's Aadhaar programme and the AI4Bharat Indian-language AI initiative at IIT Madras.

— SARVAM —

AI for all, from India.